

DANIEL GONZALEZ CEDRE

DISCRETE MATHEMATICS

UNIVERSITY OF NOTRE DAME

January 19, 2024

These notes are intended for students of CSE 20110 Discrete Mathematics at the University of Notre Dame.

Copyright © 2024 Daniel Gonzalez Cedre
<https://daniel-gonzalez-cedre.github.io>

Contents

<i>Logic</i>	6
<i>0 Language</i>	7
<i>0.1 A Brief History of...</i>	7
<i>0.2 Syntax and Semantics</i>	9
<i>0.3 A Recurring Theme</i>	10
<i>1 Zeroth-Order Logic</i>	12
<i>1.1 Propositions</i>	12
<i>1.2 Logical Connectives</i>	14
<i>1.3 The Propositional Calculus</i>	18
<i>1.4 Logical Nonequivalence</i>	24
<i>1.5 Propositional Rules of Inference</i>	25
<i>Index</i>	30

Foreword

Color	Interpretation
Yellow	<i>Emphasis</i>
Pink	<i>Definitions</i>
Purple	Internal links
Blue	External links

Table 1: Color legend.

Syntax	Transliteration	Semantics
\top	"True."	A true sentence; a tautology.
\perp	"False."	A false sentence; a contradiction.
$x := y$	" x is, by definition, y ."	The name x has been assigned to the object referenced by y .
$x = y$	" p equals q ."	p and q refer to the same object.
$p \equiv q$	" p is equivalent to q ."	The sentence p is logically equivalent to the sentence q .
$p \Leftrightarrow q$	" p if, and only if, q ."	
$p \vdash q$	" p proves q ."	By assuming the sentence p , we can prove the sentence q .
$p \Rightarrow q$	" p implies q ."	
\emptyset	"The empty set."	The set containing no elements.
$\{a, b, c\}$	"The set containing a, b , and c ."	The object whose elements are a, b, c .
$\{x \mid \varphi(x)\}$	"The set of all x such that $\varphi(x)$."	The object whose elements consist of all possible objects x for which the sentence $\varphi(x)$ is true.
$\{x \in \mathcal{A} \mid \varphi(x)\}$	"The set of all x in \mathcal{A} such that $\varphi(x)$."	The object consisting of all x from \mathcal{A} for which the sentence $\varphi(x)$ is true.
$f : \mathcal{A} \rightarrow \mathcal{B}$	" f is a function from \mathcal{A} to \mathcal{B} ."	A function named f with domain \mathcal{A} and codomain \mathcal{B} .
$f(x)$	" f of x ."	The output of f on the input x , where $x \in \mathcal{A}$ and $f(x) \in \mathcal{B}$.
\mathbb{N}	"enn"	The set of natural numbers.
\mathbb{Z}	"zee"	The set of integers.
\mathbb{Q}	"queue"	The set of rational numbers.
\mathbb{R}	"arr"	The set of real numbers.

Table 2: An overview of some important notation. Note that some expressions, like $p \equiv q$ and $p \vdash q$, have more than one equivalent notation. The middle column gives some common ways of reading each notation in English. The last column provides the meaning of each expression.

Glyph	Name	IPA
$A \alpha$	Alpha	[a]
$B \beta$	Beta	[v]
$\Gamma \gamma$	Gamma	[ɣ]
$\Delta \delta$	Delta	[ð]
$E \epsilon$	Epsilon	[e]
$Z \zeta$	Zeta	[z]
$H \eta$	Eta	[ɛ:]
$\Theta \theta$	Theta	[θ]
$I \iota$	Iota	[i:]
$K \kappa$	Kappa	[k]
$\Lambda \lambda$	Lambda	[l]
$M \mu$	Mu	[m]
$N \nu$	Nu	[n]
$\Xi \xi$	Xi	[ks]
$O \circ$	Omicron	[o]
$\Pi \pi$	Pi	[p]
$P \rho$	Rho	[r]
$\Sigma \sigma$	Sigma	[s]
$T \tau$	Tau	[t]
$Y \upsilon$	Upsilon	[y:]
$\Phi \varphi$	Phi	[f]
$X \chi$	Chi	[kʰ]
$\Psi \psi$	Psi	[ps]
$\Omega \omega$	Omega	[ɔ:]

Table 3: The Greek alphabet.

Logic

O

Language

*"No language is justly studied merely as an aid to other purposes.
It will in fact better serve other purposes, philological or
historical, when it is studied for love, for itself."*

— J. R. R. Tolkien

We communicate our thoughts to others with the use of language. This is worth reflecting on. You are probably reading this because you have some interest in computation, mathematics, logic, or are incurably bored; the goal of these notes is—in part—to provide the mathematical background necessary to study these fields at a higher level. This is particularly true for aspiring *computer scientists*, who may have some misconceptions about their field because of its misleading name,¹ and who may not be aware that the field properly and historically falls under the grand umbrella of *mathematics*.

This ambitious undertaking must therefore involve engaging with the tumultuous and violent history of mathematics. Although modern computer science is now richly interdisciplinary, the field was born during a particularly turbulent period in the late 19th and early 20th centuries AD² agitated by an existential crisis in mathematics: a crisis caused by our flagrant use of language. Here's a short summary.

0.1 A Brief History of...

The serious study of rhetoric—the art of argumentation and persuasion—as a subject in its own right dates back to at least the 5th century BC.³ Around the 3rd century BC, Euclid's 13 books of the *Elements* heralded the birth of geometry, algorithmic computation, and the first theory of numbers,⁴ where he *proved* certain statements followed from a list of *axiomatic* assumptions. This was a great achievement, establishing mathematical *proof* as a form of *argumentation* that logically deduces conclusions from a list of common assumptions. The contemporaneous Greek philosopher Theophrastus further pushed the envelope by describing the *form* these arguments and establishing their validity.



Figure 1: A fragment of book 2 from Euclid's *Elements* taken from the Oxyrhynchus papyri, dated ca. 100 AD.

¹ It's *not* about computers *nor* is it science.

² We will see later that its roots span at least to the time of Euclid in 300 BC.

³ The time of the ancient Greek sophists, who were notably opposed by Socrates, Plato, and Aristotle.

⁴ The only evidence of algorithms before this time—for multiplying, factoring, and finding square roots—dates back to Egypt and Babylon before 1600 BC.

axiom

The ancient Greeks laid the foundation for the two instrumental aspects of mathematical thought: *abstraction* and *argumentation*. Euclid abstracted what were thought to be the fundamental truths of geometry into a list of 12 *axioms*¹ so that, instead of thinking about *that* particular wall or *that* particular stick or *that* particular roof, he could make statements and observations about *quadrilaterals*, and *lines*, and *triangles* in general. These axioms were meant to encode the rules of geometry in general: the fundamental nature of what it meant to construct and measure distances, angles, and shapes. The last of these axioms would quickly become infamous.

Axiom (Parallel Postulate).

If two straight lines meet a third straight line making two interior angles that are each less than right angles, then the two lines—if they were to be extended—must intersect on that side where the interior angles are.

If you stop to think for a moment, this postulate says something very obvious. Assuming all of Euclid's other axioms, there are a few equivalent ways to restate the parallel postulate:

1. For any line L and point P not on L , there is exactly one line parallel to L passing through P .
2. The sum of angles in a triangle is 180 degrees.
3. A right triangle with side lengths A, B, C satisfies $A^2 + B^2 = C^2$.

You'll recognize this third statement as the Pythagorean *theorem*,² which is not a mere assumption!³ For the next 2000 years, the mathematical community was haunted by the thought that it was possible to *prove* the parallel postulate using the other axioms. It seemed like the rest of the axioms did such a perfectly good job of characterizing geometry that the parallel postulate *must necessarily* follow from the other axioms.

However, between 1810–1832 AD, no less than *three* papers on *hyperbolic* geometry were published, and by 1854 *Bernhardt Riemann* had developed a theory of *Riemannian* geometry on manifolds. These were all different examples of consistent models of geometry that *denied* the parallel postulate! These ideas were intensely contested: many mathematicians and natural philosophers of the time refusing to accept the notion that geometry could be non-Euclidean because *it went against their intuitive notion of how geometry should behave*.

This whole ordeal was only foreshadowing what would come at the turn of the century. In 1874, *Georg Cantor* would make a series of discoveries⁴ surrounding the nature of infinity so fundamentally opposed to common mathematical thought that he would be antagonized and ostracized for decades, causing him to suffer serious depressive crises.

¹ An *axiom* is a statement that we assume is true without justification nor proof.

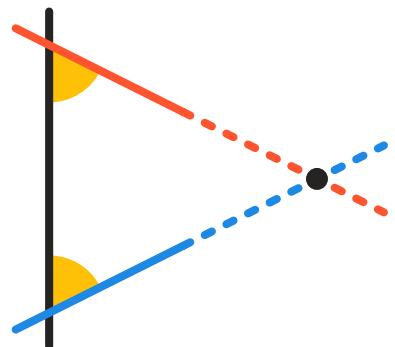


Figure 2: The parallel postulate says that any two lines — and — that make acute interior angles and with a third line — must intersect at a point .

² A *theorem* is a statement that has a proof.

³ The first two are called Playfair's axiom and the triangle postulate respectively.

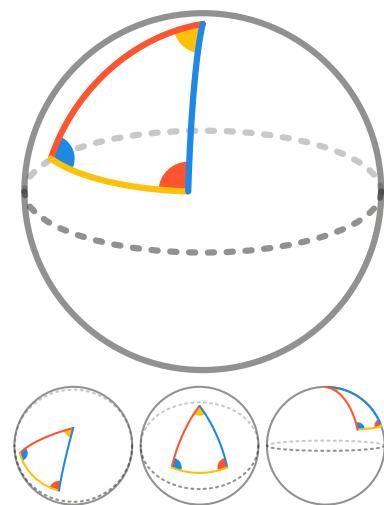


Figure 3: Four views of the same triangle whose angles sum to 270 degrees. Notice how the notions of *straight* and *parallel* differ on the surface of a sphere.

⁴ We will study these later.

Once again, mathematicians' *intuitive* notions of how *infinity* should behave were shown to be wrong. Cantor's discoveries sparked not only a civil war within the mathematical community but also a concerted effort by many mathematicians and logicians in the early 20th century to *fix* mathematics by establishing it on a firm *foundation*.

The cause of all this turmoil was, fundamentally, a *lack of precision and rigor* in the way people would communicate mathematical ideas and arguments. What does it *mean* for a line to be straight, or for two straight lines to be parallel? What does it *mean* to have two lines, or to have infinitely many lines? What *is* infinity? Is infinity a number? What *are* numbers? How do we *know* we are saying anything *true* at all?

If we hope to answer any of these questions, we must first develop a language for *precise* mathematical communication. This necessarily begins with a systematic deconstruction and analysis of *language* itself.

0.2 Syntax and Semantics

Languages encode ideas into sequences of symbols.¹ These symbols represent objects, ideas, actions, and concepts. The *meaning* behind a particular cluster of symbols is called its *semantics*. The *form* the language takes, dictated by its *grammatical rules* for composing symbols into valid sentences, is called its *syntax*. We refer to objects by giving them names. A *variable* is a symbol² that stands in place for an object that has not been determined yet.³ We can assign name to a *particular* object with the \coloneqq symbol. We call these the *terms* of an expression.

Definition 0.1 (Sentences).

A *sentence* is the expression of a complete thought or idea in accordance with the syntactic and grammatical rules of a given language. A statement is called *atomic* if it can't be broken down into smaller semantic components in any way that obeys the language's syntax and grammar.

1. A *declarative* sentence is one that describes something. They typically consist of a *subject* being described and a *predicate* property it has.
2. An *interrogative* sentence asks a non-rhetorical question.
3. An *imperative* sentence heralds a command or request.

Mathematical practice principally involves *making and justifying observations about mathematical objects*.⁴ As such, we are only really interested in crafting *declarative* sentences—sentences that describe *terms*. We will systematically deconstruct and analyse these kinds of sentences, extract their *logical essence*, and build up a new language for precisely communicating mathematical arguments.

¹ For our purposes, we will focus only on written—as opposed to spoken or signed—languages.

² We typically denote variables using single Latin or Greek letters, though there are no strict universal rules. Some common examples are listed below.

- $a, b, c, i, j, k, \ell, m, n, p, q, u, v, w, x, y, z$
- $A, B, C, D, G, H, M, N, R, X, Y, Z$
- $\alpha, \beta, \gamma, \delta, \epsilon, \eta, \theta, \lambda, \mu, \pi, \sigma, \tau, \varphi, \psi, \omega$

³ A variable does not *necessarily* refer one particular object, or even any object at all.

“Oft hope is born when all is forlorn.”

“What has it got in its pockets?”

“Keep your forked tongue behind your teeth.”

⁴ We leave the problem of *what* a mathematical object actually *is* for later.

0.3 A Recurring Theme

Before going any further, we should make a brief detour to discuss a topic that lies at the *heart* of computing, logic, and the 20th century foundational crisis in mathematics: *recursion*. In a very strong sense, what we *mean* when we say that some *thing* is *computable* is that there is a *recursive procedure* that produces that *thing*.

Idea (Church-Turing Thesis). Informally, we say something is *computable* if it can be obtained as the result of a *recursive process*, or as the output of a *Turing machine*, or expressed as a *term in the λ -calculus*.

Actually, the three concepts described above are all *equivalent* to each other. It should then be no surprise that *recursion* (and its twin *induction*) will play a central role in our studies, so we will take this brief moment to quickly describe the fundamental idea at behind recursion.¹

First, an example: how do we *compute* the sum of a list of n numbers?

$$3 + 5 + 9 + 2$$

With some hard work and determination access to the internet, we can see that $3 + 5 + 9 + 2 = 19$, but *how* did we get that answer? At the most basic level, we started by taking two of the numbers, 3 and 5 say, computing their sum $3 + 5 = 8$, and adding this intermediate result to another number from the list, 9 say, to get $8 + 9 = 17$, and adding that again to yet another element of the list—in this case, only 2 remains—to finally arrive at $17 + 2 = 19$.

$$\begin{aligned} 3 + 5 + 9 + 2 &= 3 + \cancel{5} + 9 + 2 && (1) \\ &= 8 + 9 + 2 && (2) \\ &= 8 + \cancel{9} + 2 && (3) \\ &= 17 + 2 && (4) \\ &= 17 + \cancel{2} && (5) \\ &= 19 && (6) \end{aligned}$$

This might seem so obvious it physically hurts, but let's analyse what we just did more closely. Suppose we have a list of n arbitrary numbers.²

$$x_0 + x_1 + x_2 + \cdots + x_{n-2} + x_{n-1}$$

Once again, we begin by taking the first two numbers and computing $x_0 + x_1$, then adding *this* result to x_2 , then adding *that* result to x_3 , then adding *that* result to x_4 , and so on until we reach the end of the list. So, in order to compute $x_0 + x_1 + x_2 + \dots + x_{n-2} + x_{n-1}$, we *first* need to compute $x_0 + x_1 + x_2 + \dots + x_{n-2}$ and then add that result to x_{n-1} .

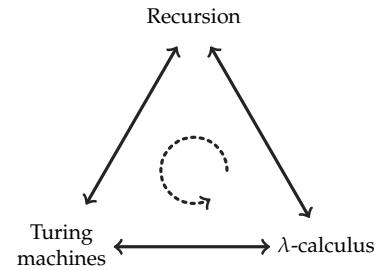


Figure 4: The Church-Turing thesis states that these three concepts—which are all *formally equivalent*—correspond with our *informal* notion of computability. In modern times, many people now take this as a definition for computability.

¹ We will leave Turing machines and the λ -calculus for a future time.

² Notice that, being the sophisticates we are, we start counting at 0, so that a list of n numbers will be indexed starting at 0 and ending at $n - 1$.

But wait, isn't $x_0 + x_1 + x_2 + \dots + x_{n-2}$ also the sum of a list? It is, it's just that the list has one less element! So *how do we compute the sum of elements in a list?* We first *compute the sum of elements in a list*, and then add one more element to that result. So, it seems like in order to do what we want, we need to already know how to do what we want; the key here is that we only need to know how to sum the elements of a *smaller* list in order to get the result we want for the *larger* list. As long as we can *eventually* get a result for one of these "*smaller*" sums, we will be able to build up a solution to our original problem by passing this result "*back up*" the chain of computation. Back to our first example.

$$\begin{aligned}
 3 + 5 + 9 + 2 &= 3 + 5 + 9 + 2 && (1) \\
 &= 3 + 5 + 9 + 2 && (2) \\
 &= 3 + 5 + 9 + 2 && (3) \\
 &= 8 + 9 + 2 && (4) \\
 &= 17 + 2 && (5) \\
 &= 19 && (6)
 \end{aligned}$$

Steps (1) through (3) continually decompose the given list into sublists on the left until we have no more lists we can break up. Each one of these lists is a smaller version of the original problem, and we compute the sums of these smaller lists by breaking them down and computing *their* sublists' sums, recombining these results at the end.

This now brings us to an important point: *we can't decompose 3 any further*, because this list only has one element in it. Do we know what the sum of all numbers in a list with one element is? Of course we do: it's just *that* number. Now we can return this result *back up* to the 5 that was waiting to be added to it, and when we add them together, we can return *that* result back to the 9 that was waiting, and then return *that* result to the 2 that was waiting, finally letting us conclude that the sum over the whole list is 19. The *recurrence relation* below summarizes this.¹

$$\text{sum}(x_0, x_1, \dots, x_{n-1}) = \begin{cases} 0 & \text{if } n = 0 \\ \text{sum}(x_0, x_1, \dots, x_{n-2}) + x_{n-1} & \text{if } n \geq 1 \end{cases}$$

We've exposed here a *recurrence* and a *basis*—the two key components underlying recursion (and, later, induction). The *recurrent* part of this procedure explains how to express a problem in terms of "*smaller*" instances of the *same problem*, describing how to combine the solutions to those subproblems into a solution for the original problem. Obviously, though, if you just keep decomposing problem into subproblems forever, you'll never be able to actually generate an answer to anything. Eventually, you need to *stop* and actually say what the answer to something is. The *basis*, does exactly this by providing explicit answers to the *smallest* versions of the problem.²

This paragraph describes the *recurrence*.

This paragraph encounters the *basis*.

¹ Notice that this is actually written slightly differently than the procedure we've just described; think about *how* this is different and whether or not it actually computes the same result as the procedure we were just analysing.

recurrence
relation

recurrence

basis

² This is sometimes called the *base case*.

1

Zeroth-Order Logic

“The limits of my language means the limits of my world.”

– Ludwig Wittgenstein

As we saw in the previous chapter, sentences can be broadly classified based on the kind of information they convey—their *functional role* in language. How do we begin deconstructing the descriptive fragment of our language? Naturally, we can think to classify the descriptive sentences by asking the fundamental question: *is this description true?*

1.1 Propositions

Let's consider the following declarative sentence.

“Ahab is a captain.” (1.1)

Here we have a descriptive sentence about the term *Ahab*—a man and thus an object of our discourse—asserting he *is a captain*. In the context of Herman Melville's *Moby Dick*, this is an accurate description. Referring to the above sentence as $\sigma_{1.1}$, we would then say $\sigma_{1.1}$ is *true*. We introduce the symbol \top to denote these kinds of sentences.

“Ishmael is a whale.” (1.2)

The above sentence, however, which we will name $\sigma_{1.2}$, immediately furrows the brow and strikes at the heart of our conscience. We know from the story that Ishmael is a sailor, and thus human, and therefore *not* a whale! We should then want to say that $\sigma_{1.2}$ is *false*, reserving the symbol \perp for sentences of this kind.

The attributes *true* and *false* that we are attaching to these sentences are what we call *truth values*, and they are the essential component of the kinds of sentences we want to express. Sentences that are *true* all exhibit a quality that makes them similar to each other but dissimilar to *false* sentences, regardless what the actual sentences themselves *mean*

true
 \top

false
 \perp

truth value



Figure 1.1: Illustration by Rockwell Kent from "Moby Dick: or, The Whale."

The symbols \top and \perp are also sometimes called “*top*” and “*bot*” respectively.

semantically. What we've just done is *abstract* the fundamental concept of truth value from descriptive sentences. This abstraction allows us to notice that *all true sentences are essentially the same as each other*, at least from the perspective of their truth values, with the same applying to *false* sentences. On the other hand, *true* and *false* sentences are complete opposites. This relationship inspires our first definition below.

Definition 1.1 (Logical Equivalence).

equivalence
 \equiv

We say that two sentences φ and ψ are *logically equivalent* when they have the same truth value. We denote this by writing $\varphi \equiv \psi$.¹

With this new definition, we can formalize our observations from the preceding paragraph as $\sigma_{1.1} \equiv \top$ and $\sigma_{1.2} \equiv \perp$ as well as $\sigma_{1.1} \not\equiv \sigma_{1.2}$. Notice that each of these three expressions is a complete sentence describing properties² held by some objects.³ In fact, these statements were themselves *true* declarative sentences. Now, let's ponder the following sentence, which we will call $\sigma_{1.3}$.

"Colorless green ideas sleep furiously." (1.3)

Like the previous examples, this is a grammatically correct, declarative sentence, but what does this sentence *mean*? Is it *true*? Is it *false*? Taking the normal English definitions for each of the words in this sentence, it doesn't seem to make any sense. We then clearly can't call it an accurate description of anything, so it can't possibly be *true*. Does that mean it must be *false*? Well, if we assume it is *false*, then what about the following sentence?

"Colorless green ideas *do not* sleep furiously." (1.4)

nonsense

This one, which we will call $\sigma_{1.4}$, seems to be saying the opposite of whatever $\sigma_{1.3}$ was saying, so if the other one is *false*, then this one must be *true*. The question then becomes: what is $\sigma_{1.4}$ accurately describing? *This* sentence seems to make just as little sense as the original! This should lead us to conclude that $\sigma_{1.3}$ could not have been *false* either, so that sentence *has no truth value!* We call expressions like this *nonsensical* because they *carry no semantic meaning*.

Let's now analyse the following statement, which we will call $\sigma_{1.5}$.

"This sentence is false." (1.5)

Expressed a little more *formally*, this is the sentence—named $\sigma_{1.5}$ —that says $\sigma_{1.5} \equiv \perp$. This certainly doesn't seem like nonsense; it says something clear about a well-understood object. So, what is the truth value of this sentence? We can try reasoning about this like we did before by examining the two possible truth values the $\sigma_{1.5}$ can take.

¹ " φ is logically equivalent to ψ ."

² being (or not) logically equivalent

³ the sentences $\sigma_{1.1}$ and $\sigma_{1.2}$

First, let's assume $\sigma_{1.5}$ is true, which we write formally as $\sigma_{1.5} \equiv \top$. By definition, this would imply $\sigma_{1.5}$ is an accurate description of some object, so we should believe what the sentence says about that object. In this case, the object is $\sigma_{1.5}$ and the description is that $\sigma_{1.5} \equiv \perp$. This *contradicts* our initial assumption! ↴ Therefore, $\sigma_{1.5}$ is *not true*¹

That rules out one truth value. What happens then if we assume $\sigma_{1.5}$ is false? Again, we can write this formally as $\sigma_{1.5} \equiv \perp$. By definition, this implies we should *reject* what $\sigma_{1.5}$ is asserting, leaving us with $\sigma_{1.5} \neq \perp$. As before, a *contradiction* emerges! ↴ Therefore, $\sigma_{1.5}$ is *not false* either!

From this simple analysis, we can see that $\sigma_{1.5}$ *does not have a truth value*. Sentences that *contradict themselves* like this are called *paradoxes*.² In the preceding analysis, we relied on the idea that \top and \perp are opposed to each other, so that the same sentence can't meaningfully be both \top and \perp at the same time. This should be intuitive based on our natural understanding and usage of the words *true* and *false*, but we will make it a point to *formally* introduce this idea now.

Axiom (Principle of Bivalence).

Any sentence expressing a truth value is either *true* or *false* but not both.

What this analysis has hopefully shown us is that *not every* well-formed, declarative sentence expresses a truth value. In order for a sentence to express a truth value, it must satisfy the following three properties.

1. The sentence must be grammatically well-formed.
2. The sentence must be declarative.
3. The sentence must be semantically meaningful.

These are the kinds of statements are *eligible to carry a truth value*—the ones for which *it would make sense* to say they are either *true* or *false*—so they will form the foundation of our new language. We will eventually call these *propositions*, but beware that this is not (yet) a *formal* definition of what a proposition is. First, we need to get a better sense of *what* propositions are linguistically and *how* they are formed.

¹ We conclude this because this is the opposite of our initial assumption, which lead us to a contradiction.

² The word *paradox* is unfortunately overload and context-dependent. When referring to specific sentences, we will use it to specifically mean a self-contradictory sentence such as $\sigma_{1.5}$, but it is also commonly used in some contexts to refer to situations that are simply *unintuitive* rather than outright contradictory.

1.2 Logical Connectives

The examples of sentences we've seen so far have all been *atomic*—meaning they can't be broken down into simpler sentences that themselves are complete thoughts—but we can obviously express thoughts that are more than merely atomic. These *compounded* propositions are formed by taking smaller propositional sentences and *connecting* them together based on what our intended meaning is.

p	q	$\neg p$	$p \wedge q$	$p \vee q$	$p \rightarrow q$	$p \leftrightarrow q$
T	T	⊥	T	T	T	T
T	⊥	⊥	⊥	T	⊥	⊥
⊥	T	T	⊥	T	T	⊥
⊥	⊥	T	⊥	⊥	T	T

Table 1.1: A truth table summarizing the basic connectives of classical logic. The two left-most columns represent the *input* values of the propositions p and q . The remaining columns describe the *output* of each expression given the corresponding inputs on each row.

Each of these different ways of connecting sentences together suggests a different way of *transforming* between truth values by combining the truth values of the component propositions into a truth value for the compound expression.

In this section, we will uncover these different transformations—which we will call *logical connectives*—and encode them using *truth tables*, which specify the output truth values for every combination of inputs.

logical
connective

Negations

Suppose we encountered the following sentence, which we call $\sigma_{1.6}$.

$$\text{"Espresso is not delicious."} \quad (1.6)$$

Immediately, the moral observer will realize the offensive absurdity of this sentence, compelled by the force of conscience to declare $\sigma_{1.6} \equiv \perp$! With this, we could simply carry on with our day; however, pausing to think for a moment, we can see that $\sigma_{1.6}$ is intimately related to the following (much more pleasant) sentence, which we call $\sigma_{1.7}$.

$$\text{"Espresso is delicious."} \quad (1.7)$$

negation

This sentence is *clearly true*, letting us sigh $\sigma_{1.7} \equiv \top$ in relief. Not only that, it is the saying exactly the opposite of what $\sigma_{1.6}$ asserted! We call propositions like these *negations* of each other. This is our first example of a *transformation* of truth value: the negation of a proposition is another proposition with the opposite truth value. To denote this formally, we introduce the \neg symbol, allowing us to write $\sigma_{1.6} \equiv \neg\sigma_{1.7}$.

We can now think of \neg formally as a *unary function* that operates on truth values.¹ This function works by mapping $\neg\top$ to \perp and by mapping $\neg\perp$ to \top . This gives us a way of abstracting negations at the level of truth values, so that we can formally define what it means to *negate* a proposition. We provide this definition now in Table 1.2, where the left-most columns² represent the inputs to \neg and the right-most columns show the truth value of the resulting output expression.³

Table 1.2: Truth table for negations.

p	$\neg p$
T	⊥
⊥	T

¹ A function is *unary* if it takes only one input argument. We will study functions in more detail later.

² shown with white backgrounds

³ shown with colored backgrounds

Conjunctions & Disjunctions

But we can obviously connect two (and sometimes more) sentences together to create larger sentences in English. For example,

“Espresso is delicious, and it nourishes the soul.” (1.8)

This sentence is composed of two smaller atomic sentences, namely “espresso is delicious” and “espresso nourishes the soul,” which we know are both independently *true*. Connecting them together with the word “and” should then, based on the way this word works in English, produce another *true* sentence. Conversely, if either of the subexpressions had been *false*, the compound result should also be *false*. This *binary* connective is called the logical *conjunction*, and we denote it using the \wedge symbol. It is defined in [Table 1.3](#).

There are several distinct ways this connective can appear in English that are nonetheless equivalent. Some examples are listed below.

-
- “Espresso is delicious, *and* it nourishes the soul.”
 - “Espresso is delicious *and* soul-nourishing.”
 - “Espresso is delicious, *but* it nourishes the soul.”
 - “Espresso is delicious, *yet* nourishing to the soul.”
 - “Espresso is delicious; *further*, it nourishes the soul.”
 - “*Although* espresso is delicious, it *also* nourishes the soul.”
-

The conjunction has a *logical dual* called the *disjunction*, defined in [Table 1.3](#) using the \vee symbol and exemplified by the following sentence.

“Espresso is delicious, or it nourishes the soul.” (1.9)

We all these two connectives *dual* to each other because negating all of the inputs to one of them is equivalent to negating the output of the other. Specifically, the expression $\neg p \vee \neg q$ is equivalent to $\neg(p \wedge q)$ whenever p and q are propositions.

Definition 1.2 (Logical Duality).

We say two logical connectives f and g are *logically dual* if negating the inputs of f is always logically equivalent to negating the output of g . Equivalently, we can say f is *logically dual* to g if applying f after \neg always gives the same result as applying \neg after g on any given inputs.

Conjunctions and disjunctions are just one example of a dual connective pair. In fact, every logical connective is dual to some other connective!¹ For now, we present this result about \wedge and \vee *without proof*; we will prove this statement when we discuss [Theorem 1.5](#) in a short while.

Table 1.3: Truth table for logical conjunctions and disjunctions.

p	q	$p \wedge q$	$p \vee q$
T	T	T	T
T	⊥	⊥	T
⊥	T	⊥	T
⊥	⊥	⊥	⊥

Table 1.4: These sentences are all logically equivalent to σ_{1.8}, though this list is obviously not exhaustive.

¹ Why might this be? Think about this.

Conditional Statements

We turn our attention now to sentence $\sigma_{1.10}$ below.

"If espresso nourishes the soul, then I will drink it." (1.10)

This is a *conditional* sentence, composed of two subclauses called the *antecedent* and the *consequent*.¹ When we use this sort of linguistic construction, we mean to say that *if* the premise happens, *then* the conclusion must also happen. Said another way: the conclusion must occur *whenever* the premise is satisfied. Notice *we are not asserting anything* about the antecedent or consequent individually! We are only establishing a *relationship* where the consequent occurs *every time* that the premise is satisfied. We call this the *material implication*, denoted by the \rightarrow symbol and defined in Table 1.5.

$p_{1.10} :=$ "Espresso nourishes the soul."

$q_{1.10} :=$ "I will drink espresso."

The antecedent and consequent for $\sigma_{1.10}$ are defined above. With these definitions, we can now write $\sigma_{1.10} \equiv p_{1.10} \rightarrow q_{1.10}$ and observe that $\sigma_{1.10}$ simply says: *if* $p_{1.10} \equiv \top$, *then* $q_{1.10} \equiv \top$. Importantly, this is *the only thing* that $\sigma_{1.10}$ is asserting! This sentence *is not saying* that if $p_{1.10} \equiv \perp$, then $q_{1.10} \equiv \perp$. In fact, if the premise is *false*, then $\sigma_{1.10}$ says *nothing* about whether or not $q_{1.10}$ is *true* or *false*.

To make this concrete, suppose I told you the following.

"If you make an A in this class, then I will eat my shoe." (1.11)

If you do happen to make an A in this class, then I'll be forced to physically eat my shoe in order to keep up my end of the bargain; in that case, the sentence was true.² On the other hand, if you make an A- instead, then I can go home with both shoes and conscience intact; in this case, the sentence was also true.³ However, what if you make the A- but I decide to eat my shoe anyways? Did I lie? No; just because you failed to make an A doesn't mean I *can't* eat my shoe! All I said was that I definitely would if you made an A.⁴ That sentence is only a lie when you *do* make an A in the class, but I refuse to eat my shoe, since I'm then breaking my promise.⁵

"I will drink espresso *if* it nourishes the soul."

"Espresso nourishes the soul *only if* I drink it."

"It is *sufficient* that espresso nourish the soul for me to drink it."

"It is *necessary* that I drink espresso for it to nourish the soul."

"I will drink espresso *unless* it doesn't nourish the soul."

Table 1.5: Truth table for conditionals.

p	q	$p \rightarrow q$	$p \leftrightarrow q$
\top	\top	\top	\top
\top	\perp	\perp	\perp
\perp	\top	\top	\perp
\perp	\perp	\top	\top

¹ Synonyms for *antecedent* & *consequent*.

protasis	apodosis
sufficient	necessary
premise	inference
assumption	conclusion
supposition	deduction
implicant	implicand
hypothesis	thesis

² $\top \rightarrow \top \equiv \top$

³ $\perp \rightarrow \perp \equiv \top$

⁴ $\perp \rightarrow \top \equiv \top$

⁵ $\top \rightarrow \perp \equiv \perp$

Table 1.6: These sentences are *all* logically equivalent to $\sigma_{1.10}$. Pay close attention to grammar of each sentence, and make special note of *where* the connectives appear.

biconditional Finally, the *material equivalence*,¹ also called the *biconditional* and written $p \leftrightarrow q$, is *true* exactly when p and q have the same truth value and is *false* otherwise. With these connectives all defined, we are now ready to formally introduce the *recursive definition* of a proposition.

¹ This is often written “if and only if” in English, abbreviated *iff*.

1.3 The Propositional Calculus

Definition 1.3 (Proposition).

proposition We say that λ is a *proposition* iff λ satisfies the following recurrence.

1. $\lambda = \top$ or $\lambda = \perp$.
2. $\lambda = \neg(\varphi)$, where φ is a proposition.
3. $\lambda = (\varphi) \wedge (\psi)$ where φ and ψ are propositions.
4. $\lambda = (\varphi) \vee (\psi)$, where φ and ψ are propositions.
5. $\lambda = (\varphi) \rightarrow (\psi)$ where φ and ψ are propositions.
6. $\lambda = (\varphi) \leftrightarrow (\psi)$, where φ and ψ are propositions.

This definition works by first establishing the *base case* that \top and \perp are propositions. We then *inductively* define larger propositions by composing together already-existing propositions with logical connectives. Initially, the only known propositions are \top and \perp . We can then take those values and build propositions like $\neg\perp$ and $(\top) \wedge (\perp)$. By continually iterating, we can see that complicated expressions like $((\neg\top) \wedge (\perp)) \rightarrow \top$ satisfy our definition of a proposition.

There is a small, but important, caveat here: what about expressions like

Definition 1.4 (Propositional Formula).

A *propositional formula* is an expression that evaluates as a proposition when all of its *variables* are themselves replaced by propositions.

Logical Equivalence

p	q	$\neg(p \wedge q)$	$\neg p \vee \neg q$
\top	\top	\perp	\top
\top	\perp	\top	\top
\perp	\top	\top	\top
\perp	\perp	\top	\top

Notice the use of *equality* = rather than *equivalence* \equiv throughout this definition. In each statement here, we are saying that the statement λ is *equal* to the expression on the right-hand side of the = symbol, meaning *they are the same sentence written in the same way*. This gives a *syntactic* definition of what a proposition is.

The use of parentheses in this definition is to avoid issues with order of operations; in situations where the meaning is clear, we can *carefully* drop parentheses.

The astute reader may have noticed that some expressions are logically equivalent to each other even if they are symbolized in different ways. For example, it's clear that $p \leftrightarrow q \equiv (p \rightarrow q) \wedge (q \rightarrow p)$, as the name “if

Table 1.7: A truth table verifying that the propositions $\neg(p \wedge q)$ and $\neg p \vee \neg q$ are logically equivalent.

and only if" suggests. We saw another example of an equivalence when we examined the duality of \wedge and \vee in ?? . We can see that statements like these are logically equivalent because, whenever we make the same assignment of truth values to the variables in these expressions, the resulting truth value is always the same. This amounts to writing out a joint truth table for the two expressions and observing that the output columns are identical.

p	q	$p \rightarrow q$	$\neg q \rightarrow \neg p$
T	T	T	T
T	⊥	⊥	⊥
⊥	T	T	T
⊥	⊥	T	T

Table 1.8: A truth table illustrating the logical equivalence between $p \rightarrow q$ and its *contrapositive* statement $\neg q \rightarrow \neg p$.

Taking this approach means having to construct a truth table whenever we want to check whether or not two propositions are equivalent. This would obviously get tedious very quickly if we were doing all of our work by hand, but it would be straightforward to automate this with only a few lines of code. If we are given two propositions $\varphi(p_1, \dots, p_n)$ and $\psi(p_1, \dots, p_n)$ that each consist of the propositional variables p_1 through p_n , then all we have to do to answer $\varphi(p_1, \dots, p_n) \stackrel{?}{=} \psi(p_1, \dots, p_n)$ is go through all possible combinations of truth assignments to p_1 through p_n and check that the outputs of $\varphi(p_1, \dots, p_n)$ and $\psi(p_1, \dots, p_n)$ match. Now, p_1 takes on one of two truth values: either T or ⊥. For each of these, we then have check both truth values p_2 can take, and two more truth values each for p_3 , and so on.

If $n = 1$, that means we only need two rows in our truth table to exhaust the entire search space. But any time we introduce a new variable to an expression, we *double* the size of the search space because we now have *two more truth values* to check for *each* row in the truth table we had already built. We can express this relationship *recursively* as follows.

$$\text{rows}(n) = \begin{cases} 1 & \text{if } n = 0 \\ 2 \cdot \text{rows}(n - 1) & \text{if } n \geq 1 \end{cases} \quad (1.12)$$

This clearly means we have to construct a truth table with 2^n rows in order to answer our question. This obviously doesn't scale computationally; it quickly becomes impossible to even *allocate enough space* for our output columns, much less compute and check them.

proof
Instead, we can *prove* the statement by constructing a formal logical argument that derives the conclusion from the premises. This involves starting from a commonly-agreed-upon list of assumptions—called *axioms*—and applying *rules of inference* until the conclusion is deduced.

axiom

Axioms of Propositional Logic

Before we talk about the rules of inference, we should first take a look at the *foundational assumptions* we will make about this logical system. In [Table 1.9](#), we have the axioms for propositional logic.

Name	Conjunctive Fragment	Disjunctive Fragment
Identity	$p \wedge \top \equiv p$	$p \vee \perp \equiv p$
Complement	$p \wedge \neg p \equiv \perp$	$p \vee \neg p \equiv \top$
Commutativity	$p \wedge q \equiv q \wedge p$	$p \vee q \equiv q \vee p$
Associativity	$p \wedge (q \wedge r) \equiv (p \wedge q) \wedge r$	$p \vee (q \vee r) \equiv (p \vee q) \vee r$
Distributivity	$p \wedge (q \vee r) \equiv (p \wedge q) \vee (p \wedge r)$	$p \vee (q \wedge r) \equiv (p \vee q) \wedge (p \vee r)$
Conditional Disintegration		$p \rightarrow q \equiv \neg p \vee q$
Biconditional Disintegration	$p \leftrightarrow q \equiv (p \rightarrow q) \wedge (q \rightarrow p)$	

Each of the statements in this table is a logical equivalence, and could be “verified” by constructing the appropriate truth table. The attitude we will take, however, is that each statement in the table simply *is* true *a priori*, without the need for them to actually be verified. They will be the starting point for our proofs in this section, and our attitude in-a-sense is that we should only care about universes in which these statements are true.

With that said, let’s try to prove our first *theorem*.¹ The *complement* axiom in the second row of [Table 1.9](#) shows us two important facts about the negation of any proposition. If we take a proposition p and conjoin it with its negation $\neg p$, that axiom tells us that we get \perp ; by duality, if we disjoin p with $\neg p$, then we get \top . Is this behavior *characteristic* of $\neg p$? Is it possible for another proposition, call it q , to exhibit the same behavior but *not* be equivalent to $\neg p$?

What the following theorem tells us is that *no*, that is not possible! Any proposition that behaves like the negation of p must have the same truth value as $\neg p$, so it might as well be $\neg p$ for all intents and purposes.

Theorem 1.1 (Uniqueness of Complements).

For any propositions p and q , if $p \wedge q \equiv \perp$ and $p \vee q \equiv \top$, then $q \equiv \neg p$.

Proof. Let p and q be propositions,² and assume that $p \wedge q \equiv \perp$ and $p \vee q \equiv \top$.³ We will prove this statement by showing that q and $\neg p$ are both equivalent to the same expression. First, observe that

Table 1.9: The first five axioms specify a *Boolean algebra*. The last two axioms *augment* the algebra with conditionals.

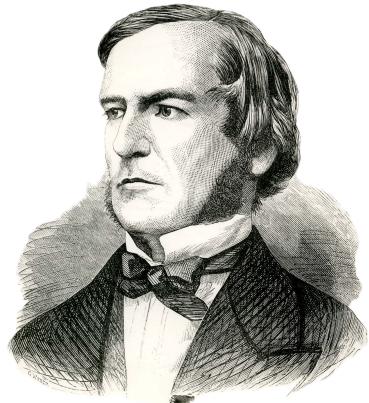


Figure 1.2: George Boole, a largely self-taught mathematician, logician, and philosopher, first described his eponymous algebras in his 1854 monograph “The Laws of Thought.”

¹ A *theorem* is a provable proposition.

² Since we need to prove this statement for any two propositions p and q , we introduce two *arbitrary* propositions at the beginning of our proof.

³ We are allowed to make these assumptions because we are proving an *if-then* statement, and they are the premise of the conditional.

$$\begin{aligned}
 \neg p &\equiv \neg p \wedge \top && \text{by } \textit{identity} \\
 &\equiv \neg p \wedge (p \vee q) && \text{because } p \vee q \equiv \top \text{ by assumption} \\
 &\equiv (\neg p \wedge p) \vee (\neg p \wedge q) && \text{by } \textit{distributivity} \\
 &\equiv \perp \vee (\neg p \wedge q) && \text{by } \textit{complement} \\
 &\equiv \neg p \wedge q && \text{by } \textit{identity}.
 \end{aligned}$$

As a result, $\neg p \equiv \neg p \wedge q$. Similarly, let's observe that

$$\begin{aligned}
 q &\equiv q \wedge \top && \text{by } \textit{identity} \\
 &\equiv q \wedge (p \vee \neg p) && \text{by } \textit{complement} \\
 &\equiv (q \wedge p) \vee (q \wedge \neg p) && \text{by } \textit{distributivity} \\
 &\equiv (p \wedge q) \vee (\neg p \wedge q) && \text{by } \textit{commutativity} \\
 &\equiv \perp \vee (\neg p \wedge q) && \text{because } p \wedge q \equiv \perp \text{ by assumption} \\
 &\equiv \neg p \wedge q && \text{by } \textit{identity}.
 \end{aligned}$$

This gives us $q \equiv \neg p \wedge q$. Therefore, we can conclude that $q \equiv \neg p$ as desired because they are both logically equivalent to $\neg p \wedge q$.

Q.E.D.

Notice how every statement in the proof above is written with *purpose*. The first sentence introduces the two *arbitrary* propositions we want to talk about. The second sentence states two assumptions (apart from the axioms in Table 1.9); we are justified in making those assumptions because we are proving a *conditional* statement. Then, we make a series of claims of logical equivalence, and every single one of those claims is *justified* by either an axiom, an assumption, or a previous theorem (which we had none of at that time).

Since we proved Theorem 1.1, we can now use this result in the future for the proofs of (often more complicated) statements. For example, it should be intuitive that $\top \equiv \neg \perp$ and $\perp \equiv \neg \top$. We can now prove this statement as a simple *corollary*¹ of Theorem 1.1.

Corollary 1.1.

$\top \equiv \neg \perp$ and $\perp \equiv \neg \top$.

Proof. Observe that $\perp \wedge \top \equiv \perp$ by the *identity* axiom. Similarly, we have that $\perp \vee \top \equiv \top \vee \perp \equiv \top$ by *commutativity* and the *identity* axiom again. So, we can apply Theorem 1.1² and conclude $\top \equiv \neg \perp$. Similarly, we can observe that $\top \wedge \perp \equiv \perp \wedge \top \equiv \perp$ by *commutativity* and *identity*, and $\top \vee \perp \equiv \perp$ by the *identity* axiom. Therefore, $\perp \equiv \neg \top$ by Theorem 1.1.

Q.E.D.

Q.E.D. stands for *Quod Erat Demonstrandum*, which is Latin for “what was to be shown has been demonstrated,” after the Greek ὃ περ ἔδει δεῖξαι. This is called a *tombstone*, and it is a traditional way of denoting the end of a proof. Modern authors might use \square or \blacksquare instead.

¹ An immediate consequence of a theorem that follows with very little effort is called a *corollary*.

² We can invoke the theorem here because we have just *proven* the premises of the theorem are true for the particular propositions we are looking at (in this case, $p := \perp$ and $q := \top$). That means, since we've satisfied the premise of the theorem, that we get to assert the conclusion, and we can cite the theorem as our justification for that assertion.

A proof gives us more than just a formal verification of a statement; it also tells us that the statement is a *necessary consequence* of the axioms we assumed in setting up our logical system. This means that we didn't have to explicitly *assume* or *define* \top to be opposite to \perp because this fact is true in *any* Boolean algebra.

Another intuitive fact that is easily taken for granted is that logical equivalence is preserved under negations. This allows us to take two equivalent expressions, negate "both sides" of the equivalence, and still obtain equivalent expressions.

Corollary 1.2.

For any propositions p and q , if $p \equiv q$, then $\neg p \equiv \neg q$.

Proof. Let p and q be propositions such that $p \equiv q$. Observe that

$$\begin{aligned} q \wedge \neg p &\equiv p \wedge \neg p && \text{because we assumed } p \equiv q \\ &\equiv \perp && \text{by complement.} \end{aligned}$$

Similarly, we can see

$$\begin{aligned} q \vee \neg p &\equiv p \vee \neg p && \text{because we assumed } p \equiv q \\ &\equiv \top && \text{by complement.} \end{aligned}$$

Therefore, we can apply [Theorem 1.1](#) and conclude that $\neg p \equiv \neg q$.

Q.E.D.

Corollary 1.3.

For any propositions p, q, r, s such that $p \equiv q$ and $r \equiv s$, we have

$$\begin{aligned} p \wedge r &\equiv q \wedge s \\ p \vee r &\equiv q \vee s \\ p \rightarrow r &\equiv q \rightarrow s \\ p \leftrightarrow r &\equiv q \leftrightarrow s \end{aligned}$$

We include these facts here just for completeness, so that we have the basic properties of \equiv codified somewhere.

Our goal for the rest of this section will be to prove [De Morgan's laws](#). Towards this goal, let's set up a few more important facts.

Theorem 1.2 (Double Negation).

For any proposition p , we have that $p \equiv \neg\neg p$.

Proof. Let p be a proposition. We will show that $p \equiv \neg\neg p$ by showing that p acts like the negation of $\neg p$. Observe that $\neg p \wedge p \equiv p \wedge \neg p \equiv \perp$ by *commutativity* and the *complement* axiom. We can similarly see, again by *commutativity* and *complement*, that $\neg p \vee p \equiv p \vee \neg p \equiv \top$. Therefore, by [Theorem 1.1](#), we can conclude that $p \equiv \neg(\neg p)$.

As an aside: another way to have gone about this section would have been to first show [Theorem 1.2](#) and [Corollary 1.2](#), then approach [Corollary 1.1](#) by first showing $\top \equiv \perp$ then citing [Corollary 1.2](#) to get $\neg\top \equiv \neg\perp$, and then citing [Theorem 1.2](#) to conclude $\neg\top \equiv \perp$.

Q.E.D.

Theorem 1.3 (Idempotency).

For any proposition p , we have $p \wedge p \equiv p$ and $p \vee p \equiv p$.

Proof. Let p be a proposition and observe that

$$\begin{aligned} p \wedge p &\equiv (p \wedge p) \vee \perp && \text{by identity} \\ &\equiv (p \wedge p) \vee (p \wedge \neg p) && \text{by complement} \\ &\equiv p \wedge (p \vee \neg p) && \text{by distributivity} \\ &\equiv p \wedge \top && \text{by complement} \\ &\equiv p && \text{by identity.} \end{aligned}$$

Similarly,

$$\begin{aligned} p \vee p &\equiv (p \vee p) \wedge \top && \text{by identity} \\ &\equiv (p \vee p) \wedge (p \vee \neg p) && \text{by complement} \\ &\equiv p \vee (p \wedge \neg p) && \text{by distributivity} \\ &\equiv p \vee \perp && \text{by complement} \\ &\equiv p && \text{by identity.} \end{aligned}$$

Therefore, we have $p \wedge p \equiv p$ and $p \vee p \equiv p$ as desired.

Q.E.D.

Theorem 1.4 (Domination).

For any proposition p , we have $p \vee \top \equiv \top$ and $p \wedge \perp \equiv \perp$.

Proof. Let p be a proposition and observe that

$$\begin{aligned} p \vee \top &\equiv p \vee (p \vee \neg p) && \text{by complement} \\ &\equiv (p \vee p) \vee \neg p && \text{by associativity} \\ &\equiv p \vee \neg p && \text{by idempotency} \\ &\equiv \top && \text{by complement.} \end{aligned}$$

Similarly,

$$\begin{aligned} p \wedge \perp &\equiv p \wedge (p \wedge \neg p) && \text{by complement} \\ &\equiv (p \wedge p) \wedge \neg p && \text{by complement} \\ &\equiv p \wedge \neg p && \text{by idempotency} \\ &\equiv \perp && \text{by complement} \end{aligned}$$

Therefore, can conclude that $p \vee \top \equiv \top$ and $p \wedge \perp \equiv \perp$.

Q.E.D.

We now finally have all the tools we need. Let's use the results we've learned so far to prove [De Morgan's laws](#).

Theorem 1.5 (De Morgan's Laws).

$\neg(p \wedge q) \equiv \neg p \vee \neg q$ and $\neg(p \vee q) \equiv \neg p \wedge \neg q$ for all propositions p, q ,

Proof. Let p and q be propositions. We will prove this theorem by showing $\neg p \vee \neg q$ acts like the negation of $p \wedge q$ and then applying *uniqueness of negations*. Observe

$$\begin{aligned}
 (p \wedge q) \wedge (\neg p \vee \neg q) &\equiv p \wedge (q \wedge (\neg p \vee \neg q)) && \text{by associativity} \\
 &\equiv p \wedge ((q \wedge \neg p) \vee (q \wedge \neg q)) && \text{by distributivity} \\
 &\equiv p \wedge ((q \wedge \neg p) \vee \perp) && \text{by complement} \\
 &\equiv p \wedge (q \wedge \neg p) && \text{by identity} \\
 &\equiv p \wedge (\neg p \wedge q) && \text{by commutativity} \\
 &\equiv (p \wedge \neg p) \wedge q && \text{by associativity} \\
 &\equiv \perp \wedge q && \text{by complement} \\
 &\equiv \perp && \text{by domination}
 \end{aligned}$$

Further, we can see clearly that

$$\begin{aligned}
 (p \wedge q) \vee (\neg p \vee \neg q) &\equiv ((p \wedge q) \vee \neg p) \vee \neg q && \text{by associativity} \\
 &\equiv (\neg p \vee (p \wedge q)) \vee \neg q && \text{by commutativity} \\
 &\equiv ((\neg p \vee p) \wedge (\neg p \vee q)) \vee \neg q && \text{by distributivity} \\
 &\equiv ((p \vee \neg p) \wedge (\neg p \vee q)) \vee \neg q && \text{by commutativity} \\
 &\equiv (\top \wedge (\neg p \vee q)) \vee \neg q && \text{by complement} \\
 &\equiv (\neg p \vee q) \vee \neg q && \text{by identity} \\
 &\equiv \neg p \vee (q \vee \neg q) && \text{by associativity} \\
 &\equiv \neg p \vee \top && \text{by complement} \\
 &\equiv \top && \text{by domination}
 \end{aligned}$$

Therefore, we can conclude $\neg(p \wedge q) \equiv \neg p \vee \neg q$ by [Theorem 1.1](#).

Q.E.D.

1.4 Logical Nonequivalence

Showing that two propositional expressions are *not* equivalent is considerably simpler than showing that they *are*. As we saw, in order to check that two propositional expressions are equivalent, we either need a proof, or we need to compute an exponentially-sized truth table. This is because, by definition, two expressions are logically equivalent when they have the same truth value *for any* given assignment of truth values to the variables in the expressions.¹ So, in order for this to *not* be the case, we only need that the expressions should disagree in truth value *for some particular* truth assignment. Instead of an entire truth table, all we need is a *single row* in which the output columns disagree.

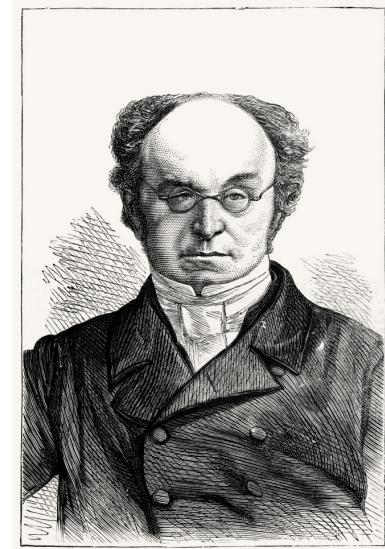


Figure 1.3: [Augustus De Morgan](#), after whom the famous laws are named.

Table 1.10: Truth table showing the nonequivalence $\neg(p \wedge q)$ and $\neg p \wedge \neg q$.

p	q	$\neg(p \wedge q)$	$\neg p \wedge \neg q$
T	T	⊥	⊥
T	⊥	T	⊥
⊥	T	T	⊥
⊥	⊥	T	T

¹ cf., subsection 1.3, first paragraph.

non
equivalence
 \neq

Definition 1.5 (Logical Nonequivalence).

Let φ and ψ be propositional expressions both consisting of the *same* variables p_1, \dots, p_n . We say that φ is *not equivalent* to ψ if there is an assignment of truth values to the variables making the truth value of φ different from the truth value of ψ . In this case, we write $\varphi \neq \psi$.

For example, to show that $p \rightarrow q \neq q \rightarrow p$, all we have to do is let $p := \top$ and $q := \perp$. We can then observe that $p \rightarrow q \equiv \top \rightarrow \perp \equiv \perp$, while $q \rightarrow p \equiv \perp \rightarrow \top \equiv \top$. Thus, we can conclude $p \rightarrow q \neq q \rightarrow p$, and this would be a full and complete proof.

1.5 Propositional Rules of Inference

Name	Inference Rule	Interpretation
Modus Ponens	$p, (p \rightarrow q) \vdash q$	If we have $p \rightarrow q$ and we know p , then we can deduce q .
Modus Tollens	$\neg q, (p \rightarrow q) \vdash \neg p$	If we have $p \rightarrow q$ but also $\neg q$, then we can infer $\neg p$.
Reductio ad Absurdum	$(\neg p \rightarrow q), (\neg p \rightarrow \neg q) \vdash p$	If, by assuming $\neg p$, we are able to deduce q and $\neg q$, then $\neg p$ is absurd and we must have p instead.
The Deduction Theorem	$(p \vdash q) \vdash (p \rightarrow q)$	If, by assuming p , we can construct a proof of q , then we can write $p \rightarrow q$.
Truth Axiom	$\vdash \top$	We can always write \top without any prior assumptions.

So far, we've developed a modestly-powerful formal language capable of expressing some basic logical ideas founded on *axioms* that give us a syntactic grounding for our semantic interpretation of the language. The axioms in Table 1.9—establishing that the logic of propositions behaves algebraically like a Boolean algebra—are all *equivalences*. They give us truth-preserving ways of manipulating expressions, telling us that when we replace certain syntactic patterns with other ones, we don't affect the resulting truth value. We've now used these equivalences in the proofs of some basic facts about our language, and we know we are justified in the use of these equivalences in our proofs because they form the axiomatic foundation of our formal system. However, our proofs do not consist *only* of equivalences. In fact, every one of our proofs has necessarily relied on some form of reasoning that is more than just an equivalence statement saying that one expression can be syntactically replaced by another.

This is most apparent in the proof of Theorem 1.1, our very first theorem. We began that proof by introducing two arbitrary propositions and then *immediately assuming* that their conjunction was \perp and their disjunction as \top . Those assumptions were not based on any of the equivalence

Table 1.11: The rules of inference for the propositional logic.

rules of
inference

axioms, like *complement* or *identity*, that involve \perp and \top , so why were we allowed to say that in our proof? Further, in the proof of [Corollary 1.1](#), we apply this theorem by saying that, since we've satisfied the premises of this theorem for a pair of particular propositions, we get to write down the conclusion of the theorem. Why does it make sense to say that? Nothing yet in our language *formally* gives us the ability to do these things, but we have been doing them in our proofs because they are natural, logical ways of expressing certain kinds of arguments. In order to codify this in our formal language, we need to introduce this kind of reasoning—which engages in *inferences* rather than equivalences—in the form of more axioms for our language. Importantly, however, these axioms will work in *one direction only*, from left to right, unlike equivalences, which work symmetrically. We call these inference axioms the *rules of inference* for classical logic.

+
prove

The rules in [Table 1.11](#) each take the form $\Gamma \vdash \varphi$,¹ where Γ represents a list of assumptions and φ is the conclusion that follows from them. The \vdash symbol, sometimes called a turnstile, expresses the idea that we can *prove* φ on the basis of Γ (in addition to the rest of our axioms and any other theorems we have shown). In the event that Γ is empty—as in the Truth Axiom in our table of rules of inference—this simply means that the conclusion φ can be inferred *without* any assumptions.

The first, and most important, of the rules of inference is *modus ponens*,² enabling us to *build up* arguments that conclude in a point. *Modus ponens* is, in a sense, the *essence* of argumentation and rhetoric. Without it, the conclusion of a *conditional* statement would not be *conditioned* on its premise, and the words “premise” and “conclusion” would suddenly lose their meaning. There would be no point in establishing an argument of any sort about anything because the argument would never actually lead anywhere. Its sister rule, *modus tollens*,³ conversely allows us to *break down* arguments by denying premises when their promised conclusions don’t occur.

The third rule, *reductio ad absurdum*,⁴ gives us the ability to construct proofs by contradiction. Suppose we are interested in proving some proposition p . If, hypothetically, we had assumed $\neg p$, and from there had been able to infer that both q and $\neg q$ must be the case, where q is some other proposition, then that must mean that $\neg p$ induces a *contradiction*. Hence, $\neg p$ can’t possibly be true, yielding p as a result. This form of argumentation—like arguing with a hammer—is incredibly powerful and has been in use since at least the year 400 BC.⁵

The last two rules of inferences are technical necessities for us to interact meaningfully with our formal language. The *deduction theorem* tells us that proofs correspond to syntactic implications in our language. When

¹ This can be read as “ Γ proves φ ” or as “by assuming Γ we can deduce φ ” in English.

² *Modus ponens* is short for the Latin phrase *modus ponendo ponens*, literally “the method of putting by placing.”

³ *Modus tollens* is short for the Latin phrase *modus tollendo tollens*, literally “the method of removing by taking away.”

⁴ *Reductio ad absurdum* is a Latin phrase meaning “reduction to the absurd.” This is also called *argumentum ad absurdum*, denoting an “argument to absurdity.”

⁵ In Plato’s dialogues, Socrates frequently engages in this sort of reasoning by showing his opponents’ seemingly-sensible statements can be systematically dismantled to absurdity.

we *prove* that $p \vdash q$ by first *assuming p* and then *deriving q*, that actually means that we also have a proof of $p \rightarrow q$. The *truth axiom* asserts we can always introduce the sentence \top in our proof at any time without having to satisfy any prior conditions.

Although this is a rather small collection of rules, it is capable of representing any kind of expressible propositional rhetoric. Despite this, it's not a *minimal* set of axioms for the zeroth-order logic. In fact, it's possible to have an even smaller set of rules without sacrificing the rhetorical strength of our language. *Modus tollens*, for instance, could actually be shown to follow from the other rules of inference as a *theorem*, reducing our total number of assumptions. Let's prove it now.

Theorem 1.6 (Modus Tollens).

We have $\neg q, (p \rightarrow q) \vdash \neg p$ for any propositions p and q .

Proof. Let p and q be arbitrary propositions, and suppose $\neg q$ and also $p \rightarrow q$. We know that $p \rightarrow q \equiv \neg q \rightarrow \neg p$,¹ so we have $\neg q \rightarrow \neg p$. Then, by *modus ponens*, we can conclude $\neg p$.

Q.E.D.

¹ This result—that a conditional statement is equivalent to its *contrapositive*—is an exercise in your problem set.

We don't provide a *minimal* set of inference rules in order to keep this topic somewhat manageable, but the interested reader might be excited to learn that *all* of propositional logic can be encoded using *just two connectives* (\neg and \rightarrow) and *just three axioms* along with *modus ponens*.

In addition to these rules, there are several classical *syllogisms*²—studied since times of ancient Greece—that have formed the historical bedrock for argumentation in philosophy. We can prove some of these now.

Theorem 1.7 (Conditional Elimination).

We have $(p \rightarrow q) \vdash (p \vdash q)$ for any propositions p and q .

Proof. Let p and q be arbitrary propositions, and suppose $p \rightarrow q$. We will now show that $p \vdash q$. Assume p . Then, since we have $p \rightarrow q$, we can derive q by *modus ponens*. Thus, $p \vdash q$.

Q.E.D.

² A *syllogism* is a traditional way to refer to a deductive argument that draws a conclusion from a small set of premises.

This theorem, when combined with the *deduction theorem*, establishes the syntactic equivalence between the semantically-distinct \rightarrow and \vdash symbols.

Theorem 1.8 (Hypothetical Syllogism).

We have $(p \rightarrow q), (q \rightarrow r) \vdash p \rightarrow r$ for any propositions p , q , and r .

Proof. Let p , q , and r be arbitrary propositions, and suppose $p \rightarrow q$ and $q \rightarrow r$. We will first show that $p \vdash r$. Assume p . Since $p \rightarrow q$, we have q by *modus ponens*. Further, since we have $q \rightarrow r$, we get r by *modus ponens*. Thus, $p \vdash r$. Therefore, by applying the *deduction theorem*, we can conclude $p \rightarrow r$.

Q.E.D.

This theorem allows us to construct and follow chains of reasoning, a fundamental aspect of any nontrivial argument.

Theorem 1.9 (Conjunction Introduction).

We have $p, q \vdash p \wedge q$ for any propositions p and q .

This is known as *adjunction*.

Proof. Let p and q be arbitrary propositions. Assume p , and also separately assume q . Towards a contradiction,¹ suppose $\neg(p \wedge q)$. We can plainly see

$$\begin{aligned}\neg(p \wedge q) &\equiv \neg p \vee \neg q && \text{by De Morgan's laws} \\ &\equiv p \rightarrow \neg q && \text{by conditional disintegration.}\end{aligned}$$

¹ When beginning a *proof by contradiction*, it is good form to explicitly alert the reader to this fact when you make your problematic assumption.

So, we have $p \rightarrow \neg q$, from which we can derive $\neg q$ by *modus ponens*. This shows us that $\neg(p \wedge q) \rightarrow \neg q$ by the *deduction theorem*. However, since we also had q by assumption, we can derive $\neg(p \wedge q) \rightarrow q$ using the *deduction theorem* again. \nexists ²

Therefore, we can conclude $p \wedge q$ by *reductio ad absurdum*.

Q.E.D.

² The symbol \nexists is useful for highlighting to the reader where, in a proof by contradiction, the actual *contradiction* is.

Theorem 1.10 (Conjunction Elimination).

We have $(p \wedge q) \vdash p$ for any propositions p and q .

This is known as *simplification*.

Theorem 1.11 (Disjunctive Syllogism).

We have $(p \vee q), \neg p \vdash q$ for any propositions p and q .

Theorem 1.12 (Disjunction Introduction).

We have $p \vdash (p \vee q)$ for any propositions p and q .

This is known as *addition*.

Theorem 1.13 (Disjunction Elimination).

We have $(p \rightarrow r), (q \rightarrow r), (p \vee q) \vdash r$ for any propositions p, q, r .

This is known as *proof by cases*.

Theorem 1.14 (Biconditional Elimination).

We have $(p \leftrightarrow q), p \vdash q$ for any propositions p, q .

Theorem 1.15 (Constructive Dilemma).

We have $(p \rightarrow s), (q \rightarrow t), (p \vee q) \vdash (s \vee t)$ for any p, q, s, t .

Theorem 1.16 (Ex Falso Quodlibet).

We have $p, \neg p \vdash q$ for any propositions p, q .

This is known as the *principle of explosion*.

Index

- atomic, 9
- axiom, 19
- Boolean algebra, 20
- conjunction, 16
- contrapositive, 27
- corollary, 21
- disjunction, 16
- equivalence
 - logical, 13
 - material, 18
 - nonequivalence, 25
- formula
 - propositional, 18
- logical
 - duality, 16
- material equivalence, 13
- nonequivalence, 25
- material equivalence, 18
- implication, 17
- negation, 15
- proof, 19
- proposition
 - formal, 18
- propositional formula, 18
- variable, 18
- rules of inference, 26
- sentence, 9
- theorem, 20